

# **DD2380 Artificial Intelligence**

## **Machine Learning 2: Reinforcement learning**

---

Alexandre Proutiere

September 6, 2017


KTH (The Royal Institute of Technology)

# Outline of today's lecture

- 0. Introduction: Supervised vs. Unsupervised learning
- 1. Supervised Learning
- 2. **Reinforcement Learning**
  - A. **Playing against an unknown environment**
  - B. Bandit optimisation
  - C. RL in Markov Decision Processes

# Reinforcement learning

Learning optimal sequential behaviour / control from interacting with the environment (data generated as we go in an adaptive manner)



**Unknown** state dynamics:

$$s_{t+1}^\pi = F_t(s_t^\pi, a_t^\pi)$$

# Reinforcement learning: Applications



- Making a robot walk
- Portfolio optimisation
- Playing games better than humans
- Helicopter stunt manoeuvres
- Optimal communication protocols in radio networks
- Display ads
- Search engines
- ...

# IID vs. Markovian dynamics


- IID (Independent and Identically Distributed) environment: Bandit optimisation

Each *round*, select one action and observe a random reward whose distribution depends on the action only




- Markovian environment: Learning in Markov Decision Processes (MDPs)

The new state is randomly generated depending on the former state and the selected action



# Bandit optimisation




State dynamics:

$$s_{t+1}^\pi = F_t(s_t^\pi, a_t^\pi)$$

- Interact with an i.i.d.
- The reward is independent of the state and is the only feedback:
  - i.i.d. environment:  $r_t(a, s) = r_t(a)$  random variable with mean  $\theta_a$
  - adversarial environment:  $r_t(a, s) = r_t(a)$  is arbitrary!

# Markov Decision Process



State dynamics:

$$s_{t+1}^\pi = F_t(s_t^\pi, a_t^\pi)$$

- History at  $t$ :  $h_t^\pi = (s_1^\pi, a_1^\pi, \dots, s_{t-1}^\pi, a_{t-1}^\pi, s_t^\pi)$
- Markovian environment:  $\mathbb{P}[s_{t+1}^\pi = s' | h_t^\pi, s_t^\pi = s, a_t^\pi = a] = p(s' | s, a)$
- Stationary deterministic rewards (for simplicity):  $r_t(a, s) = r(a, s)$

# What is to be learnt and optimised?

- Bandit optimisation: the average rewards of actions are unknown  
Information available at time  $t$  under  $\pi$ :

$$a_1^\pi, r_1(a_1^\pi), \dots, a_{t-1}^\pi, r_{t-1}(a_{t-1}^\pi)$$

- MDP: The state dynamics  $p(\cdot|s, a)$  and the reward function  $r(a, s)$  are unknown  
Information available at time  $t$  under  $\pi$ :


$$s_1^\pi, a_1^\pi, r_1(a_1^\pi, s_1^\pi), \dots, s_{t-1}^\pi, a_{t-1}^\pi, r_{t-1}(a_{t-1}^\pi, s_{t-1}^\pi), s_t^\pi$$

- Objective: maximise the cumulative reward

$$\sum_{t=1}^T \mathbb{E}[r_t(a_t^\pi, s_t^\pi)] \quad \text{or} \quad \sum_{t=1}^{\infty} \lambda^t \mathbb{E}[r_t(a_t^\pi, s_t^\pi)]$$



# Regret




- Difference between the cumulative reward of an "Oracle" policy and that of agent  $\pi$
- Regret quantifies the price to pay for learning!
- Exploration vs. exploitation trade-off: we need to probe all actions to play the best later ...

# Outline of today's lecture

- 0. Introduction: Supervised vs. Unsupervised learning
- 1. Supervised Learning
- 2. **Reinforcement Learning**
  - A. Playing against an unknown environment
  - B. **Bandit optimisation**
  - C. RL in Markov Decision Processes


# Bandit Optimisation

**First application:** Clinical trial, Thompson 1933




- A set of possible actions at each step
- Unknown sequence of rewards for each action
- Bandit feedback: only rewards of chosen actions are observed
- Goal: maximise the cumulative reward (up to step  $T$ )

## Rate adaptation in 802.11 wireless systems



- The AP sequentially sends packets to the receiver and has  $K$  available encoding rates  $r_1 < r_2 < \dots < r_K$  + MIMO modes
- The unknown probability a packet sent at rate  $r_k$  is received is  $\theta_k$
- Goal: design a rate selection scheme that learns the  $\theta_k$ 's and quickly converges to rate  $r_{k^*}$  maximising  $\mu_k = r_k \theta_k$  over  $k$

## Search engines




- The engine should list relevant webpages depending on the request 'jaguar'
- The CTRs (Click-Through-Rate) are unknown
- Goal: design a list selection scheme that learns the list maximising its global CTRs

# Stochastic bandit taxonomy

**Unstructured problems:** average rewards are not related

$$\theta = (\theta_1, \dots, \theta_K) \in \Theta = \prod_k [a_k, b_k]$$

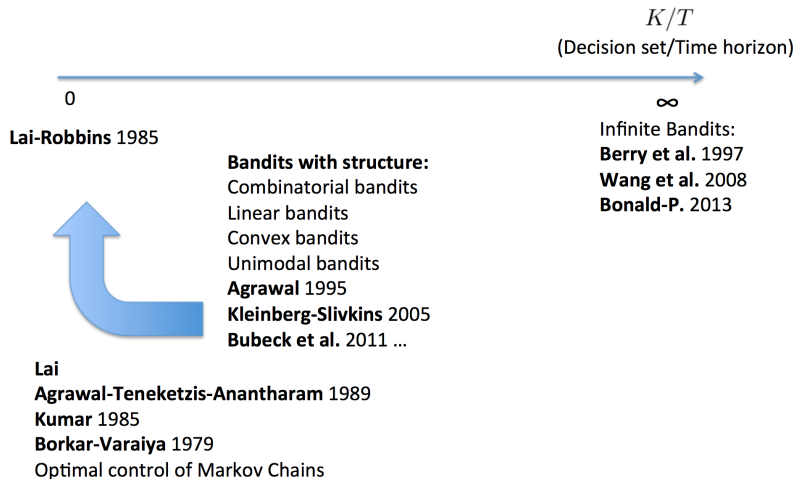


**Structured problems:** the decision maker knows that average rewards are related. She knows  $\Theta$ . The rewards observed for a given arm provides side information about the other arms.

$\theta = (\theta_1, \dots, \theta_K) \in \Theta$  not an hyper-rectangle



# Stochastic bandit taxonomy



- Discrete unstructured bandits  $A = \{1, \dots, K\}$   
Regret  $O(K \log(T))$  (stochastic),  $O(\sqrt{KT})$  (adversarial)
- Infinite Bandits  $A = \mathbb{N}$ , Bayesian setting  
Regret  $O(\sqrt{T})$
- Continuous Bandits  $A \subset \mathbb{R}^d$   
Structures:  $a \rightarrow \theta_a$  is convex, Lipschitz, linear, unimodal (quasi-convex) etc.  
Regret:  $O(\text{poly}(d)\sqrt{T})$



# Unstructured stochastic bandits – Robbins 1952

- Finite set of actions  $A = \{1, \dots, K\}$
- (Unknown) rewards of action  $a \in A$ :  $(X_t(a), t \geq 0)$  i.i.d. Bernoulli with  $\mathbb{E}[X_t(a)] = \theta_a$
- Optimal action  $a^* \in \arg \max_a \theta_a$
- Online policy  $\pi$ : select action  $a_t^\pi$  at time  $t$  depending on  $a_1^\pi, X_1(a_1^\pi), \dots, a_{t-1}^\pi, X_{t-1}(a_{t-1}^\pi)$
- Regret up to time  $T$ :  $R^\pi(T) = T\theta_{a^*} - \sum_{t=1}^T \theta_{a_t^\pi}$

# Regret lower bound

**Uniformly good algorithms:** An algorithm  $\pi$  is uniformly good if for all  $\theta \in \Theta$ , for any sub-optimal arm  $a$ , the number of times  $n_a(t)$  arm  $a$  is selected up to round  $t$  satisfies:  $\mathbb{E}[n_a(t)] = o(t^\alpha)$  for all  $\alpha > 0$ .

## Fundamental performance limits: (Lai-Robbins1985)

For any uniformly good algorithm  $\pi$ :

$$\liminf_T \frac{R^\pi(T)}{\log(T)} \geq \sum_{a \neq a^*} \frac{\theta_{a^*} - \theta_a}{KL(\theta_a, \theta_{a^*})}$$

where  $KL(a, b) = a \log(\frac{a}{b}) + (1 - a) \log(\frac{1-a}{1-b})$  (KL divergence)

. Regret linear in  $K$ , and scaling as  $1/(\theta_{a^*} - \theta_a)$

Empirical reward of arm  $a$ :  $\hat{\theta}_a(t) = \frac{1}{n_a(t)} \sum_{n=1}^t X_n(a) 1_{a(n)=a}$

- **$\epsilon$ -greedy.** In each round  $t$ :
  - with probability  $1 - \epsilon$ , select the best empirical arm  
 $a^*(t) \in \arg \max_a \hat{\theta}_a(t)$
  - with probability  $\epsilon$ , select an arm uniformly at random

The algorithm has linear regret (not uniformly good)

- **$\epsilon_t$ -greedy.** In each round  $t$ :
  - with probability  $1 - \epsilon_t$ , select the best empirical arm  
 $a^*(t) \in \arg \max_a \hat{\theta}_a(t)$
  - with probability  $\epsilon_t$ , select an arm uniformly at random

The algorithm has logarithmic regret for Bernoulli rewards and

$$\epsilon_t = \min(1, \frac{K}{t\delta^2}) \text{ where } \delta = \min_{a \neq a^*} (\theta_{a^*} - \theta_a)$$

# Algorithms

## *Optimism in front of Uncertainty*

### **Upper Confidence Bound algorithm:**

Auer et al., 2002

$$b_a(t) = \hat{\theta}_a(t) + \sqrt{\frac{2 \log(t)}{n_a(t)}}$$


$\hat{\theta}(t)$ : empirical reward of  $a$  up to  $t$

$n_a(t)$ : nb of times  $a$  played up to  $t$

In each round  $t$ , select the arm with highest index  $b_a(t)$

Under UCB, the number of times  $a \neq a^*$  is selected satisfies:

$$\mathbb{E}[n_a(T)] \leq \frac{8 \log(T)}{(\theta_{a^*} - \theta_a)^2} + \frac{\pi^2}{6}$$



**KL-UCB algorithm:** Lai 1987, Garivier et Cappe 2011

$$b_a(t) = \max\{q \leq 1 : n_a(t)KL(\hat{\theta}_a(t), q) \leq f(t)\}$$

where  $f(t) = \log(t) + 3 \log \log(t)$  is the *confidence* level.


In each round  $t$ , select the arm with highest index  $b_a(t)$

Under KL-UCB, the number of times  $a \neq a^*$  is selected satisfies: for all  $\delta < \theta_{a^*} - \theta_a$ , for all  $T$ ,

$$\mathbb{E}[n_a(T)] \leq \frac{\log(T)}{KL(\theta_a + \delta, \theta_{a^*})} + C \log \log(T) + \delta^{-2}$$


# Structured stochastic bandits

- $a \mapsto \theta_a$  is (lipschitz, convex, unimodal, multimodal, ...)
- $\theta \in \Theta$ ,  $\Theta$  encodes the known structure
- 802.11 rate adaptation example. Graphical unimodality: there is a graph (vertices = arms) such that from any vertex, there is a path to the optimal arm with increasing average reward



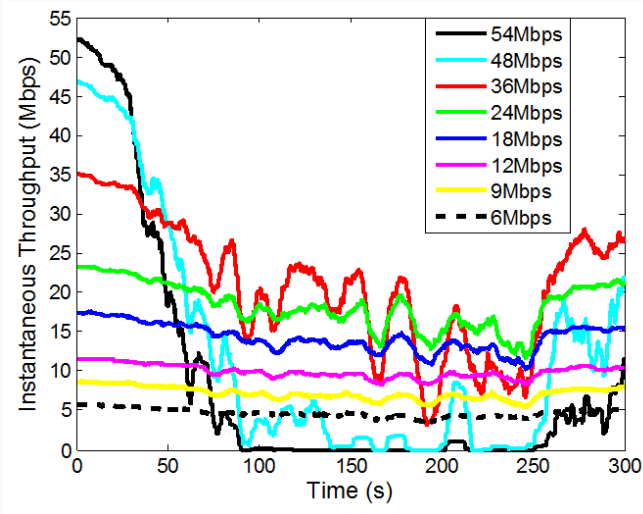
# Unimodal bandits

- Optimal exploration: Only neighbors of the optimal arm generate log regret
- Optimal algorithm<sup>1</sup>: maximises throughput over a time window



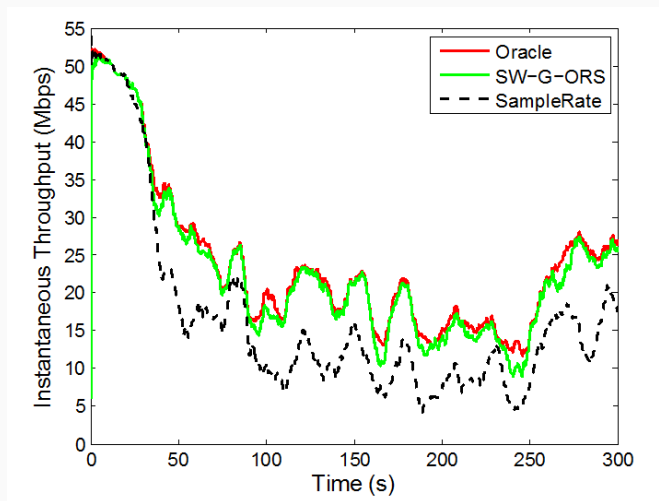
<sup>1</sup>Optimal Rate Sampling in 802.11, Combes et al., IEEE Infocom 2014

# Non-stationary environment





## SW version of optimal algorithms



# Outline of today's lecture

0. Introduction: Supervised vs. Unsupervised learning

1. Supervised Learning


2. **Reinforcement Learning**

A. Playing against an unknown environment

B. Bandit optimisation

C. **RL in Markov Decision Processes**

# Markov Decision Process (MDP)



State dynamics:

$$s_{t+1}^\pi = F_t(s_t^\pi, a_t^\pi)$$

- Markovian environment:  $\mathbb{P}[s_{t+1}^\pi = s' | h_t^\pi, s_t^\pi = s, a_t^\pi = a] = p(s' | s, a)$
- Stationary deterministic rewards (for simplicity):  $r_t(a, s) = r(a, s)$
- $p(\cdot | s, a)$  and  $r(\cdot, \cdot)$  are unknown initially

# Example

Playing pacman (Google Deepmind experiment, 2015)



**State:** the current displayed image

**Action:** right, left, down, up

**Feedback:** the score and its increments + state

# Bellman's equation

**Objective:** max the average discounted reward  $\sum_{t=1}^{\infty} \lambda^t \mathbb{E}[r(a_t^\pi, s_t^\pi)]$

Assume the transition probabilities and the reward function are known

- Value function: maps the initial state  $s$  to the corresponding maximum reward  $v(s)$
- Bellman's equation:

$$v(s) = \max_{a \in A} \left[ r(a, s) + \lambda \sum_j p(j|s, a) v(j) \right]$$

- Solve Bellman's equation. The optimal policy is given by:

$$a^*(s) = \arg \max_{a \in A} \left[ r(a, s) + \lambda \sum_j p(j|s, a) v(j) \right]$$

What if the transition probabilities and the reward function are unknown?

- Q-value function: the max expected reward starting from state  $s$  and playing action  $a$ :

$$Q(s, a) = r(a, s) + \lambda \sum_j p(j|s, a) \max_{b \in A} Q(j, b)$$

Note that:  $v(s) = \max_{a \in A} Q(s, a)$

- Algorithm: update the  $Q$ -value estimate sequentially so that it converges to the true  $Q$ -value

1. **Initialisation:** select  $Q \in \mathbb{R}^{S \times A}$  arbitrarily, and  $s_0$
2. **Q-value iteration:** at each step  $t$ , select action  $a_t$  (each state-action pair must be selected infinitely often)

Observe the new state  $s_{t+1}$  and the reward  $r(s_t, a_t)$

Update  $Q(s_t, a_t)$ :

$$Q(s_t, a_t) := Q(s_t, a_t) + \alpha_t \left[ r(s_t, a_t) + \lambda \max_{a \in A} Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

It converges to  $Q$  if  $\sum_t \alpha_t = \infty$  and  $\sum_t \alpha_t^2 < \infty$

The crawling robot ...

<https://www.youtube.com/watch?v=2iNrJx6IDeo>




# Scaling up Q-learning

Q-learning converges very slowly, especially when the state and action spaces are large ...

State-of-the-art algorithms (optimal exploration, ideas from bandit opt.):  
regret  $O(\sqrt{SAT})$

What if the action and state are continuous variables?

Example: Mountain car demo<sup>2</sup>



---

<sup>2</sup>See Sutton tutorial, NIPS 2015

# Q-learning with function approximation

**Idea:** restrict our attention to  $Q$ -value functions belonging to a family of functions  $\mathcal{Q}$

**Examples:**

1. Linear functions:  $\mathcal{Q} = \{Q_\theta, \theta \in \mathbb{R}^M\}$ ,

$$Q_\theta(s, a) = \sum_{i=1}^M \phi_i(s, a) \theta_i = \phi^\top \theta$$

where for all  $i$ ,  $\phi_i$  is linear. The  $\phi_i$ 's are linearly independent.

2. Deep networks:  $\mathcal{Q} = \{Q_{\mathbf{w}}, \mathbf{w} \in \mathbb{R}^M\}$ ,  $Q_{\mathbf{w}}(s, a)$  given as the output of a neural network with weights  $\mathbf{w}$  and inputs  $(s, a)$

# Q-learning with linear function approximation

1. **Initialisation:** select  $\theta \in \mathbb{R}^M$  arbitrarily, and  $s_0$
2. **Q-value iteration:** at each step  $t$ , select action  $a_t$  (each state-action pair must be selected infinitely often)  
Observe the new state  $s_{t+1}$  and the reward  $r(s_t, a_t)$   
Update  $\theta$ :

$$\begin{aligned}\theta &:= \theta + \alpha_t \Delta_t \nabla_{\theta} Q_{\theta}(s_t, a_t) \\ &= \theta + \alpha_t \Delta_t \phi(s_t, a_t)\end{aligned}$$

where  $\Delta_t = r(s_t, a_t) + \lambda \max_{a \in A} Q_{\theta}(s_{t+1}, a) - Q_{\theta}(s_t, a_t)$

For convergence results, see "An analysis of Reinforcement Learning with Function Approximation", Melo et al., ICML 2008

# Q-learning with function approximation

## Success stories:

- TD-Gammon (Backgammon), Tesauro 1995 (neural nets)
- Acrobatic helicopter autopilots, Ng et al. 2006
- Jeopardy, IBM Watson, 2011
- 49 atari games, pixel-level visual inputs, Google Deepmind 2015



# Questions?

Alexandre Proutiere

[alepro@kth.se](mailto:alepro@kth.se)